

Assignment_Week_9

January 20, 2021

```
[204]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.impute import SimpleImputer
```

```
[205]: df = pd.read_csv("Data01.csv")
```

```
[206]: df.shape
```

```
[206]: (400, 5)
```

```
[66]: df.head()
```

```
[66]:
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	M	\$19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000%	0
3	15603246	Female	27	57000	0
4	15804002	Male	NaN	76000	0

```
[67]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         400 non-null   int64
1   Gender          400 non-null   object
2   Age             394 non-null   object
3   EstimatedSalary 395 non-null   object
4   Purchased       400 non-null   int64
dtypes: int64(2), object(3)
memory usage: 15.8+ KB
```

```
[68]: df['Gender'].value_counts()
```

```
[68]: Female      201
      Male       191
      M          4
      F          2
      Female%    1
      Male%      1
      Name: Gender, dtype: int64
```

Clean up data series df['Gender'] using “replace” and “map” in dataframe

```
[207]: new_Gender = { "M":"Male", "F":"Female", "Female%":"Female", "Male%":"Male" }
      df = df.replace({"Gender":new_Gender})
```

```
[208]: df['Gender'].value_counts()
```

```
[208]: Female      204
      Male       196
      Name: Gender, dtype: int64
```

```
[209]: df['Gender'].unique()
```

```
[209]: array(['Male', 'Female'], dtype=object)
```

```
[210]: df['Purchased'].unique()
```

```
[210]: array([0, 1])
```

Clean up \$ in df['Age']

```
[211]: df['Age'].unique()
```

```
[211]: array(['$19 ', '35', '26', '27', nan, '32', '25', '20', '18', '29', '47',
            '45', '46', '48', '49', '31', '21', '28', '33', '30', '23', '24',
            '22', '59', '34', '39', '19', '38', '37', '42', '40', '36', '41',
            '58', '55', '52', '60', '53', '50', '56', '51', '57', '44', '43',
            '54'], dtype=object)
```

```
[212]: df['Age'] = df['Age'].str.replace('$', '')
```

```
[213]: from sklearn.impute import SimpleImputer
      imputer = SimpleImputer(missing_values=np.NaN, strategy='mean')
      df['Age'] = imputer.fit_transform(df['Age'].values.reshape(-1,1))
```

```
[214]: df['Age'].unique()
```

```
[214]: array([19.      , 35.      , 26.      , 27.      , 37.58629442,
            32.      , 25.      , 20.      , 18.      , 29.      ,
            47.      , 45.      , 46.      , 48.      , 49.      ,
```

```

31.          , 21.          , 28.          , 33.          , 30.          ,
23.          , 24.          , 22.          , 59.          , 34.          ,
39.          , 38.          , 37.          , 42.          , 40.          ,
36.          , 41.          , 58.          , 55.          , 52.          ,
60.          , 53.          , 50.          , 56.          , 51.          ,
57.          , 44.          , 43.          , 54.          ])
```

```
[215]: df['Age'] = df['Age'].astype('int64')
```

```
[216]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       395 non-null   object
4   Purchased             400 non-null   int64
dtypes: int64(3), object(2)
memory usage: 15.8+ KB
```

```
[217]: df['Age'].unique()
```

```
[217]: array([19, 35, 26, 27, 37, 32, 25, 20, 18, 29, 47, 45, 46, 48, 49, 31, 21,
        28, 33, 30, 23, 24, 22, 59, 34, 39, 38, 42, 40, 36, 41, 58, 55, 52,
        60, 53, 50, 56, 51, 57, 44, 43, 54])
```

```
[218]: df['EstimatedSalary'].unique()
```

```
[218]: array(['19000', '20000', '43000%', '57000', '76000', '58000', '84000',
        '150000', '33000', '65000', '$80,000 ', '52000', '86000', '18000',
        '82000', '80000', '25000', '26000', '28000', '29000', '22000',
        '49000', '41000', '23000', '30000', '43000', '$18,000 ', '74000',
        '137000', '16000', '44000', '90000', '27000', '72000', '31000',
        '17000', '51000', '108000', '15000', '79000', '54000', '135000',
        '89000', '32000', '83000', '55000', '48000', '117000', '87000',
        '66000', '120000', '63000', '68000', '113000', '112000', '42000',
        '88000', '62000', '118000', '85000', '81000', '50000', '116000',
        '$15,000 ', '123000', '73000', '37000', '59000', '149000', '21000',
        '35000', '71000', '61000', '75000', '53000', '107000', '96000',
        '45000', '47000', '100000', '38000', '69000', '148000', '115000',
        '34000', '60000', '70000', '36000', '39000', '134000', '101000',
        '130000', '114000', '142000', nan, '78000', '143000', '91000',
        '144000', '102000', '126000', '133000', '147000', '104000',
```

```
'146000', '$122,000 ', '97000', '95000', '131000', '77000',
'125000', '106000', '141000', '93000', '138000', '119000',
'122000', '105000', '99000', '129000', '46000', '64000', '139000'],
dtype=object)
```

```
[219]: df['EstimatedSalary']=df['EstimatedSalary'].str.replace("$","")
df['EstimatedSalary']=df['EstimatedSalary'].str.replace(",","")
df['EstimatedSalary']=df['EstimatedSalary'].str.replace("%","")
```

```
[220]: imputer = SimpleImputer(missing_values=np.NaN, strategy='mean')
df['EstimatedSalary'] = imputer.fit_transform(df['EstimatedSalary'].values.
↳reshape(-1,1))
```

```
[221]: df['EstimatedSalary'].unique()
```

```
[221]: array([ 19000.      ,  20000.      ,  43000.      ,  57000.      ,
        76000.      ,  58000.      ,  84000.      , 150000.      ,
        33000.      ,  65000.      ,  80000.      ,  52000.      ,
        86000.      ,  18000.      ,  82000.      ,  25000.      ,
        26000.      ,  28000.      ,  29000.      ,  22000.      ,
        49000.      ,  41000.      ,  23000.      ,  30000.      ,
        74000.      , 137000.      ,  16000.      ,  44000.      ,
        90000.      ,  27000.      ,  72000.      ,  31000.      ,
        17000.      ,  51000.      , 108000.      ,  15000.      ,
        79000.      ,  54000.      , 135000.      ,  89000.      ,
        32000.      ,  83000.      ,  55000.      ,  48000.      ,
       117000.      ,  87000.      ,  66000.      , 120000.      ,
        63000.      ,  68000.      , 113000.      , 112000.      ,
        42000.      ,  88000.      ,  62000.      , 118000.      ,
        85000.      ,  81000.      ,  50000.      , 116000.      ,
       123000.      ,  73000.      ,  37000.      ,  59000.      ,
       149000.      ,  21000.      ,  35000.      ,  71000.      ,
        61000.      ,  75000.      ,  53000.      , 107000.      ,
        96000.      ,  45000.      ,  47000.      , 100000.      ,
        38000.      ,  69000.      , 148000.      , 115000.      ,
        34000.      ,  60000.      ,  70000.      ,  36000.      ,
        39000.      , 134000.      , 101000.      , 130000.      ,
       114000.      , 142000.      ,  70017.72151899,  78000.      ,
       143000.      ,  91000.      , 144000.      , 102000.      ,
       126000.      , 133000.      , 147000.      , 104000.      ,
       146000.      , 122000.      ,  97000.      ,  95000.      ,
       131000.      ,  77000.      , 125000.      , 106000.      ,
       141000.      ,  93000.      , 138000.      , 119000.      ,
       105000.      ,  99000.      , 129000.      ,  46000.      ,
        64000.      , 139000.      ])
```

```
[222]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID                400 non-null   int64
1   Gender                 400 non-null   object
2   Age                    400 non-null   int64
3   EstimatedSalary        400 non-null   float64
4   Purchased              400 non-null   int64
dtypes: float64(1), int64(3), object(1)
memory usage: 15.8+ KB

```

```
[223]: df.describe()
```

```

[223]:
count      User ID      Age  EstimatedSalary  Purchased
count  4.000000e+02  400.000000      400.000000  400.000000
mean    1.569154e+07  37.577500    70017.721519  0.357500
std     7.165832e+04  10.336882    33977.615953  0.479864
min     1.556669e+07  18.000000    15000.000000  0.000000
25%     1.562676e+07  30.000000    44000.000000  0.000000
50%     1.569434e+07  37.000000    70017.721519  0.000000
75%     1.575036e+07  45.250000    88000.000000  1.000000
max     1.581524e+07  60.000000   150000.000000  1.000000

```

```
[231]: df.columns
```

```

[231]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased'],
dtype='object')

```

```
[232]: df = df.rename(columns={'Age': 'age'})
```

Split data into X the predictors dataframe and y the target series

```

[233]: X = df.iloc[:, :-1]
       y = df.iloc[:, -1]

```

```

[234]: print(X.shape)
       print(y.shape)

```

```

(400, 4)
(400,)

```

```

[235]: print(type(X))
       print(type(y))

```

```

<class 'pandas.core.frame.DataFrame'>
<class 'pandas.core.series.Series'>

```

```
[236]: X.head()
```

```
[236]:   User ID  Gender  age  EstimatedSalary
0  15624510   Male   19         19000.0
1  15810944   Male   35         20000.0
2  15668575  Female   26         43000.0
3  15603246  Female   27         57000.0
4  15804002   Male   37         76000.0
```

```
[237]: y.head()
```

```
[237]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: Purchased, dtype: int64
```

Find the correlation between the three predictors ['Age', 'EstimatedSalary', 'Gender'] and the target 'Purchased'

```
[238]: corr1 = np.corrcoef(df['age'], df['Purchased'])
      corr1
```

```
[238]: array([[1.          , 0.61158189],
      [0.61158189, 1.          ]])
```

```
[239]: corr2 = np.corrcoef(df['EstimatedSalary'], df['Purchased'])
      corr2
```

```
[239]: array([[1.          , 0.36684361],
      [0.36684361, 1.          ]])
```

How about the correlation between two numerical predictors 'Age' and 'EstimatedSalary'


We want to check whether “multicollinearity” is a problem:

```
[241]: corr2 = np.corrcoef(df['EstimatedSalary'], df['age'])
      corr2
```


```
[241]: array([[1.          , 0.1694131],
      [0.1694131, 1.          ]])
```

EDA using Visualization


```
[242]: sns.scatterplot(data=df, x="age", y="Purchased")
      plt.show()
```




```
[243]: sns.scatterplot(data=df, x="EstimatedSalary", y="Purchased")  
plt.show()
```




```
[244]: ax = sns.boxplot(x="Purchased", y="age", data=df)
```




```
[245]: ax = sns.boxplot(x="Purchased", y="EstimatedSalary", data=df)
```





```
[246]: ax = sns.boxplot(x="Purchased", y="age", hue="Gender", data=df, palette="Set3")
```



```
[247]: sns.histplot(data=df, x="age", hue="Purchased", element="step", kde=True)  
plt.show()
```



```
[248]: sns.histplot(data=df, x="EstimatedSalary",  
                  hue="Purchased", element="step", kde=True)  
plt.show()
```



Predictive Models using (1) Logistic Regression and (2) Random Forest

The goal is to predict the “Purchase” behavior using “Age”, “EstimatedSalary”, and “Gender”:

```
[249]: X.head()
```

```
[249]:
```

	User ID	Gender	age	EstimatedSalary
0	15624510	Male	19	19000.0
1	15810944	Male	35	20000.0
2	15668575	Female	26	43000.0
3	15603246	Female	27	57000.0
4	15804002	Male	37	76000.0

```
[250]: y.head()
```

```
[250]:
```

0	0
1	0
2	0
3	0
4	0

Name: Purchased, dtype: int64

Convert String value ‘Gender’ into dummy variables

```
[251]: import pandas as pd  
import numpy as np
```

```
[252]: X = pd.get_dummies(X, columns=["Gender"])
```

```
[253]: X.head()
```

```
[253]:
```

	User ID	age	EstimatedSalary	Gender_Female	Gender_Male
0	15624510	19	19000.0	0	1
1	15810944	35	20000.0	0	1
2	15668575	26	43000.0	1	0
3	15603246	27	57000.0	1	0
4	15804002	37	76000.0	0	1

```
[254]: X.drop('User ID',axis=1,inplace=True)  
X.head()
```

```
[254]:
```

	age	EstimatedSalary	Gender_Female	Gender_Male
0	19	19000.0	0	1
1	35	20000.0	0	1
2	26	43000.0	1	0

3	27	57000.0	1	0
4	37	76000.0	0	1

[255]: X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   400 non-null   int64
1   EstimatedSalary       400 non-null   float64
2   Gender_Female         400 non-null   uint8
3   Gender_Male           400 non-null   uint8
dtypes: float64(1), int64(1), uint8(2)
memory usage: 7.2 KB
```

[256]: X['Gender_Female'] = X['Gender_Female'].astype('int64')

[257]: X['Gender_Male'] = X['Gender_Male'].astype('int64')

[258]: X.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 4 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   400 non-null   int64
1   EstimatedSalary       400 non-null   float64
2   Gender_Female         400 non-null   int64
3   Gender_Male           400 non-null   int64
dtypes: float64(1), int64(3)
memory usage: 12.6 KB
```

[260]: `from sklearn.linear_model import LogisticRegression`
`from sklearn.ensemble import RandomForestClassifier`
`from sklearn.model_selection import train_test_split`
`from sklearn.metrics import accuracy_score`

[261]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=0)
logreg = LogisticRegression()
logreg.fit(X_train, y_train)

[261]: LogisticRegression()

```
[262]: y_pred = logreg.predict(X_test)
print(f'Accuracy of logistic regression classifier on test set:␣
      ↳{accuracy_score(y_test, y_pred)*100:.2f}%')
```

Accuracy of logistic regression classifier on test set: 70.00%

Random Forest Classification

```
[263]: clf=RandomForestClassifier(n_estimators=100)
      clf.fit(X_train,y_train)
```

```
[263]: RandomForestClassifier()
```

```
[265]: y_pred=clf.predict(X_test)
print(f'Accuracy of Random Forest classifier on test set:␣
      ↳{accuracy_score(y_test, y_pred)*100:.2f}%')
```

Accuracy of Random Forest classifier on test set: 91.25%

Feature Importance:

```
[267]: print(clf.feature_importances_)
      print(X.columns)
```

```
[0.48984002 0.49518898 0.00739093 0.00758006]
```

```
Index(['age', 'EstimatedSalary', 'Gender_Female', 'Gender_Male'],
      dtype='object')
```

```
[269]: # The feature importance sum up to 1
import pandas as pd
feature_imp = pd.Series(clf.feature_importances_, index=X.columns).
      ↳sort_values(ascending=False)
feature_imp
```


```
[269]: EstimatedSalary    0.495189
age                    0.489840
Gender_Male           0.007580
Gender_Female        0.007391
dtype: float64
```

Feature Importance in Visualization format

```
[271]: import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

# Creating a bar plot
sns.barplot(x=feature_imp, y=feature_imp.index)
# Add labels to your graph
```

```
plt.xlabel('Feature Importance Score')
plt.ylabel('Features')
plt.title("Visualizing Important Features")
plt.show()
```



Hyperparameter Tuning Random Forest: (1) number of decision trees and (2) number of predictors for splitting trees

The default values for (1) # of trees is “100” and (2) number of predictors is Square Root

```
[293]: # Grid searching key hyperparameters for RandomForestClassifier
from sklearn.model_selection import RepeatedStratifiedKFold
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier

# define models and parameters
model = RandomForestClassifier()
n_estimators = [10, 100, 150, 200, 250, 300]
max_features = ['sqrt', 'log2', 'None']

# define grid search
grid = dict(n_estimators=n_estimators,max_features=max_features)
cv = RepeatedStratifiedKFold(n_splits=5, n_repeats=3, random_state=1)
```

```

grid_search = GridSearchCV(estimator=model, param_grid=grid, n_jobs=-1, cv=cv,
    ↳scoring='accuracy', error_score=0)
grid_result = grid_search.fit(X, y)

# summarize results
print(f"Best Accuracy: {grid_result.best_score_*100: 0.2f}% with parameters:
    ↳{grid_result.best_params_}")

```

Best Accuracy: 88.17% with parameters: {'max_features': 'sqrt', 'n_estimators': 200}

```

[294]: # Print out the performance, in Accuracy, of other combinations:
means = list(grid_result.cv_results_['mean_test_score'])
stds = list(grid_result.cv_results_['std_test_score'])
params = list(grid_result.cv_results_['params'])

# Zip function will zip iterables, in this case 3 lists, into tuple
# and then perform "tuple unpacking" into components:
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

0.876667 (0.036477) with: {'max_features': 'sqrt', 'n_estimators': 10}
0.879167 (0.037823) with: {'max_features': 'sqrt', 'n_estimators': 100}
0.879167 (0.034055) with: {'max_features': 'sqrt', 'n_estimators': 150}
0.881667 (0.038424) with: {'max_features': 'sqrt', 'n_estimators': 200}
0.878333 (0.038586) with: {'max_features': 'sqrt', 'n_estimators': 250}
0.879167 (0.039441) with: {'max_features': 'sqrt', 'n_estimators': 300}
0.873333 (0.033187) with: {'max_features': 'log2', 'n_estimators': 10}
0.877500 (0.036856) with: {'max_features': 'log2', 'n_estimators': 100}
0.877500 (0.038243) with: {'max_features': 'log2', 'n_estimators': 150}
0.875833 (0.039651) with: {'max_features': 'log2', 'n_estimators': 200}
0.880000 (0.038406) with: {'max_features': 'log2', 'n_estimators': 250}
0.877500 (0.039843) with: {'max_features': 'log2', 'n_estimators': 300}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 10}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 100}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 150}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 200}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 250}
0.000000 (0.000000) with: {'max_features': 'None', 'n_estimators': 300}

```

0.1 Let's try out the XGBoost model that won many Kaggle competition!

```

[296]: !pip install xgboost

```

Defaulting to user installation because normal site-packages is not writeable
 Looking in indexes:
<http://nexus.opr.statefarm.org/repository/pypi.python.org/simple>,
<http://nexus.opr.statefarm.org/repository/python-internal/simple>

Collecting xgboost

Downloading http://nexus.opr.statefarm.org/repository/pypi.python.org/packages/xgboost/1.3.3/xgboost-1.3.3-py3-none-manylinux2010_x86_64.whl (157.5 MB)

| 157.5 MB 5.8 MB/s eta 0:00:0101

Requirement already satisfied: numpy in /opt/conda/lib/python3.8/site-packages (from xgboost) (1.19.4)

Requirement already satisfied: scipy in /opt/conda/lib/python3.8/site-packages (from xgboost) (1.5.3)

Installing collected packages: xgboost

Successfully installed xgboost-1.3.3

```
[297]: from xgboost import XGBClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
```

```
[300]: # split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=2)
```

```
[301]: print(X_train.shape)
print(X_test.shape)
```

(320, 4)

(80, 4)

```
[302]: # fit model no training data
model = XGBClassifier()
model.fit(X_train, y_train)
```

/users/yycs/.local/lib/python3.8/site-packages/xgboost/sklearn.py:888:

UserWarning: The use of label encoder in XGBClassifier is deprecated and will be removed in a future release. To remove this warning, do the following: 1) Pass option use_label_encoder=False when constructing XGBClassifier object; and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ..., [num_class - 1].

warnings.warn(label_encoder_deprecation_msg, UserWarning)

[17:40:53] WARNING: ../src/learner.cc:1061: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

```
[302]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
importance_type='gain', interaction_constraints='',
learning_rate=0.300000012, max_delta_step=0, max_depth=6,
min_child_weight=1, missing=nan, monotone_constraints='()',
n_estimators=100, n_jobs=88, num_parallel_tree=1, random_state=0,
```



```
reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,  
tree_method='exact', validate_parameters=1, verbosity=None)
```

```
[305]: # make predictions for test data  
y_pred = model.predict(X_test)  
accuracy = accuracy_score(y_test, y_pred)  
print("Accuracy: %.2f%%" % (accuracy * 100.0))
```

Accuracy: 86.25%

0.2 The performance of XGBoost (with default setting) is not better than Random Forest (Hyperparameter tuned).

0.3 The next thing to try is to tune its long list of hyperparameters.

```
[ ]:
```